

Extensible Math Functions for C++

Document #: P4188R1
Date: 2026-03-30
Project: Programming Language C++
Audience: SG6
SG18
SG20
LEWG
Reply-to: Stéphane Gros-Lemesre
<stephane.groslemesre@gmail.com>

Abstract

This paper proposes making standard library mathematical functions extensible to user-defined types via a member function or ADL. To preserve backward compatibility, the proposed approach introduces a new `std::math` sub-namespace. The new namespace also offers an opportunity to introduce a C++-idiomatic math library consistent with general standard library naming patterns.

Contents

1	Motivation	3
1.1	Problem	3
1.1.1	Naive call sites	3
1.1.2	The <code>using</code> -statement workaround	3
1.1.3	Standard Library numeric types	5
1.2	High-level proposal	6
1.3	Use Cases	6
1.3.1	Semantic Types	6
1.3.2	Arbitrary types	6
1.3.3	Code usage examples	6
1.4	Context	7
2	Impact on the Standard	8
2.1	<code>std::math</code>	8
2.1.1	Proposed naming	8
2.2	Bridging through opt-in	9
2.3	Scope	9
2.3.1	Future extensions	9
2.4	Summary	10
3	Design Principles	10
3.1	Two spaces, different behaviours	10
3.2	New contract	10
3.3	CPO-like customization	11
3.3.1	Dispatch	11
3.3.2	Alternative to plain CPO	11
3.4	Single header	11
3.5	Canonical function names	11
3.6	Customisation-level specifications	12
4	General Specifications for built-in arithmetic types	12
4.1	Conservative addition	12
4.2	<code>constexpr</code>	12
4.3	<code>[[nodiscard]]</code>	12

4.4	Error Handling	12
5	Technical Specification	13
5.1	Functions	13
5.1.1	Power Functions	13
5.1.2	Exponential	15
5.1.3	Trigonometric	15
5.1.4	Hyperbolic	15
5.1.5	Rounding	15
5.1.6	Sign	15
5.2	Proposed Implementation	15
5.2.1	CPO	15
5.2.2	Compatibility layer	16
5.2.3	Alternative implementation with [P2806R3] (do expressions) and [P2826R2] (replacement functions)	16
6	References	17

1 Motivation

1.1 Problem

User-defined types are central to C++ programming. The standard library inherited from the Standard Template Library its foundational principle of separating data structures and algorithms through generic programming. Containers, iterators, Customization Point Objects (CPOs), custom operators all contribute to enable a natural integration of any type, user-defined or not, in a common ecosystem.

Mathematical functions are a notable gap in this picture. Unlike operators, they are not part of overload resolution in a way that extends to user-defined types, and unlike CPOs, there is no standard hook to participate in. An author implementing an arithmetic-like type will find it easy to implement binary operations such as addition, subtraction, multiplication, and division; comparisons, hashing, or integration with containers for their user-defined type, until they face the problem of implementing the square-root operation or similar mathematical function.

The fundamental problem with defining mathematical functions for a custom type is that because there are no well-defined contract like for the operators, and no CPO in the standard library for that purpose, the mechanism for their integration is not in the hands of the author. A custom implementation can be offered, but whether it will be picked up by generic code depends entirely on how that code was implemented.

This is not what good separation of concerns looks like: for the author of numeric types, providing custom mathematical functions is a leap of faith, while for generic code authors, supporting custom types might not be the primary concern and requires a significant amount of boilerplate.

1.1.1 Naive call sites

It is difficult to quantify the amount of generic code calling mathematical functions without support for user-defined numeric types. This is made even more difficult as these functions can live in the global namespace if `<math.h>` is included. Using GitHub search functionality shows 115k C++ files using the idiom `using std::pow;` for 3.5M C++ files calling `pow()`, of which 741k C++ files calling explicitly `std::pow()`.

Although these numbers have to be taken with caution, as there are other ways to leverage ADL for the `pow` function than `using std::pow;` specifically, and some of the qualified `std::pow` calls might come from non-generic code, it seems likely that a non-negligible amount of generic code does not allow custom types because of this extensibility limitation.

The library `nholthaus/units`, for instance, provides `units::math::sqrt` which calls `std::sqrt` directly on the underlying scalar value, which means it does not enable ADL resolution and thus does not extend to custom underlying types. [\[nholthaus-units\]](#)

This should not be surprising: the responsibility of making custom types work intuitively should belong to their authors, not to their users or third-party. It takes conscious effort for a generic library author to anticipate the needs for wrappers of hypothetical custom types and explicitly provide support for them.

Yet, the consequence of this status quo is substantial: the successful integration of a custom-type in a code-base no longer depends on its own design and implementation, but on the implementation of every other piece of code that will manipulate it. It is effectively a dependency of that type on all of the generic code it interacts and will interact with in the future. Such open-ended, long-term consequences are often unacceptable.

This is significant because custom numeric types can be used to support **correctness and safety** by encoding constraints that are enforced by the type system automatically. This pattern is similar to *newtypes* in Rust, and is sometimes referred to as “semantic-typing”. Whether it is by preventing inappropriate conversion or ensuring unit consistency, it can prevent entire classes of bugs and is a well-understood tool for writing safer numeric code. As it stands, its practicality in C++ is significantly restrained by how well such custom types compose with the broader ecosystem, and mathematical functions are a substantial weak point in that composition.

1.1.2 The using-statement workaround

The idiomatic workaround is to enable ADL via a `using` declaration:

```
using std::sqrt;
return sqrt(x);
```

This allows a user-defined mathematical function in the same namespace as the type passed into the function to be found via ADL, while falling back to the `std`-defined version for built-in types. However, this idiom has a critical limitation: it requires a statement, and is therefore unavailable in contexts that only accept expressions.

1.1.2.1 Statement-only

Constructor member initializer lists:

```
MyType(A aSq, B b, C c)
: a(sqrt(aSq)), // std::sqrt not found through conversion
  b(b),
  c(abs(c))      // std::abs not found through conversion
{}
```

The same applies to default member initializers:

```
struct Foo {
    double x = sqrt(v); // std::sqrt not found through conversion
};
```

requires expressions:

```
template<typename T>
concept numeric = requires(T x) {
    { abs(x) }; // requirement fails
};
```

There are other, less obvious situations where the same limitation applies, such as constant expressions (array size from a new-type) and template arguments.

In all of these cases, the programmer faces the same three unsatisfactory choices:

Option 1: Call `std::` explicitly

```
MyType(A aSq, B b, C c)
: a(std::sqrt(aSq)), // silently breaks extensibility
  b(b),
  c(std::abs(c))
{}
```

This compiles and works for built-in types, but silently cuts off any user-defined overload.

Option 2: Restructure to allow a statement

For member initializer lists, this means moving initialization to the constructor body:

```
MyType(A aSq, B b, C c)
: b(b)
{
    using namespace std;
    this->a = sqrt(aSq); // requires a to be default-constructible
    this->c = abs(c);     // loses const and reference member support
}
```

This restores ADL but members can no longer be `const` or references, and all members must be default-constructible. Not all contexts can be restructured this way.

Option 3: Write boilerplate helper functions

```
template<typename T> auto my_sqrt(T&& x) {
    using std::sqrt;
    return sqrt(std::forward<T>(x));
}

MyType(A aSq, B b, C c)
    : a(my_sqrt(aSq)),
      b(b),
      c(my_abs(c))
{}
```

This restores extensibility but requires every author of generic code to write and maintain their own dispatch layer; one wrapper per math function, 45+ at least, and likely more in the future (see [\[P3935R1\]](#)).

The ownership of these wrappers is unclear. Custom numeric-type authors should probably embed them in their math function implementations, but authors of generic code using such functions cannot rely on this being provided and may need to implement them redundantly to support a wider range of types.

1.1.2.2 Existing Practice

The existence of multiple widely-used C++ libraries implementing exactly this machinery is evidence that there are real use-cases and that the status quo is encouraging unnecessary duplication.

mp-units, **Eigen**, and **Boost.Units** have all independently converged on the same core mechanism: bring `std::sqrt` into scope and let ADL do the work.

```
using std::sqrt;
return sqrt(x);
```

The surrounding machinery differs:

- **mp-units** adds an explicit member function check. [\[mp-units\]](#)
- **Eigen** wraps the dispatch in a traits struct with SIMD specialisations and relies on macros to minimize the duplication between different functions. [\[eigen-math\]](#)

```
#define EIGEN_MATHFUNC_IMPL(func, scalar) \
    Eigen::internal::func##_impl<typename \
    Eigen::internal::global_math_functions_filtering_base<scalar>::type>
```

- **Boost.Units** applies the pattern to the inner value of a quantity type. [\[boost-units-sqrt\]](#)

But the fundamental approach is identical in all three.

1.1.2.3 Teachability

A difficulty with the `using` ADL idiom is its limited discoverability. Without specific guidance, it takes very specific circumstances to stumble on the problem, let alone its solution. Calling the `std` operations looks correct, compiles cleanly, and works as expected with native types. It is only when this code is used with custom types that its limitation becomes problematic.

The following exchange from [nholthaus/units](#) GitHub issue #39 [\[nholthaus-issue39\]](#) is instructive. A user reports that generic algorithms using ADL-found math functions do not work with unit types, and the library author responds:

“Honestly, I guess I just never use ADL because I pretty much exclusively use fully qualified namespaces in my code, and I didn’t put thought into it.”

In the current situation, it is easy to do the wrong thing, and difficult to do the right one.

1.1.3 Standard Library numeric types

The problem also impacts the standard library when non-arithmetic numeric types are defined, such as for `simd`.

Specifically, Mateusz Pusz, the author of the mp-units library and main contributor to the C++26 unit library [P3045R8], called out these very limitations at C++ on Sea 2025, specifically calling for a proposal to introduce CPOs for mathematical functions. [pusz-cpponseas2025]

The linear algebra facilities currently use helper functions in order to allow custom arithmetic types, and it has been acknowledged in committee that having a well defined contract to define these customization points would be very helpful.

1.2 High-level proposal

This proposal suggests introducing a new `std::math` namespace where extensible basic math functions would be defined. The extension would be made possible through a member function defined on the type, or via ADL.

Implementations would be provided for built-in types such as floating-points and for `std::complex` where appropriate.

In addition, an opt-in compatibility layer would also be added for basic math functions defined in the `std` namespace, opening a conservative and targeted bridge between legacy code calling `std::` operations directly and custom-types.

1.3 Use Cases

1.3.1 Semantic Types

Often ambiguously called “strong-typing”, we will use the term “semantic-typing” to describe the practice of preferring types that are specific and carry more semantic over general-purpose types. For instance, preferring a `GeoCoordinate` type over an `std::pair` to store a latitude and longitude.

In other languages (Haskell, Rust, Python), this is sometimes facilitated by a “newtype” feature. Newtypes are typically strictly-typed thin wrappers around existing primitive or data types for semantic-typing purposes.

C++26 Quantities and Units library [P3045R8] offers an opportunity to replace built-in types by types carrying the semantic of the unit the quantity is expressed as. E.g.: `std::quantity<std::si::metre>` instead of `double`.

Another example would be defining an index type paired with the container it is meant to index into, offering compile-time guarantees that it is not used with the wrong container.

For both these examples, changing the type of variables in existing code would work for operator usage, hashing and other common operations if the replacement type defines these operations, but it would fail as soon as a mathematical function is called without resorting to the `using`-statement workaround.

This proposal is not limited to scalar types, and would also allow defining specialisations of mathematical functions for multidimensional types such as vectors, quaternions, tensors, or matrices.

1.3.2 Arbitrary types

Any type can leverage this extension mechanism without restriction. It is for the implementers to decide whether a specific mathematical operation is meaningful for a specific type. This proposal does not define an explicit boundary in this regard, but defines the dispatching mechanism allowing the extension.

This can be likened to operator overloading where the facility for defining the operation is made available for implementers to leverage as they see fit without any contract enforcement beyond the admitted operator signatures.

1.3.3 Code usage examples

1.3.3.1 Newton p^{th} root

```
template<typename T>
T newton_pth_root(T a, T x, int p, T tolerance)
{
    T prev;
```

```

do
{
    prev = x;
    //  $f(x) = x^p - a$ ,  $f'(x) = p * x^{p-1}$ 
    x = ((p-1)*x + a / std::math::power(x, p-1)) / p;
}
while (std::math::abs(x - prev) > tolerance);
return x;
}

```

This implementation would work out of the box with `T` as `float`, `double`, `long double`, but also with `T` as a unit such as a velocity, or even a user-defined matrix type, assuming an implementation for `std::math::power` and `std::math::abs` are provided.

It is worth noting that this function could be implemented in the current ecosystem with one or two `using` statements: either with `using namespace std;`, or `using std::pow;` and `using std::abs;`, and then calling unqualified `pow` and `abs`, enabling ADL. The benefits of the proposed approach are perceived as providing a better-defined correct path with clearer separation of concerns (making it easier to do the right thing), and addressing the other limitation of the `using` workaround that cannot be used in expression-only contexts.

1.3.3.2 Compatibility layer example

A typical use case for the compatibility layer is a custom numeric type used with an existing library that calls `std::sqrt` directly and cannot be modified:

```

// Generic library using std::sqrt directly
namespace someLib
{
    template<typename T>
    auto someFunction(T&& someValue)
    {
        // some code...
        return std::sqrt(std::forward<T>(someValue));
    }
}

// User's custom type, with its own square root implementation in its namespace
namespace mylib
{
    struct Scalar { ... };
    Scalar square_root(Scalar x) { ... }
}

// Opt in to the compatibility layer
template<>
inline constexpr bool std::is_math_extensible<mylib::Scalar> = true;

// Now where std::sqrt(mylib::Scalar{}) is used in the library, it is
// forwarded to mylib::square_root via std::math::square_root

someLib::someFunction(mylib::Scalar{5.2}); // now works

```

1.4 Context

The mathematical functions of C++ have been inherited from C, before the introduction of namespaces in the language and therefore lived in the global name space. With C++98, the `<cmath>` header was added, as the new correct way to access the math functions inherited from C. The `<math.h>` header was retained, though, for backward compatibility, leading to the present-day situation where many mathematical functions are

accessible in both the global namespace and `std`. C++98, C++11, and C++17 also introduced additional functions to `<cmath>`. C++29 could add to it further [P3935R1].

As a result of this evolution, these functions are unlike other parts of the standard library:

- Their names do not follow the `snake_case` pattern.
- There are several variants of each function differentiated by prefix or suffix instead of using function overloading.
- Their error handling uses `errno` instead of exceptions or the more modern `std::expected`.

Encountering this API is surprising and confusing for new C++ learners without a C background, as the justification for these divergences are largely historical.

Another consequence of this long progression is also that there is a huge body of code relying on all of these successive evolutions. These APIs are essentially frozen, as the contracts established over the years can't reasonably be invalidated.

2 Impact on the Standard

Consequently, based on the points discussed above, this proposal aims to address some of the limitations to the extensibility of the math functions in C++ while preserving existing behaviours.

Since the core problem is the absence of a standard extension mechanism, the solution must be introduced by the standard library itself. It cannot be provided by user code or third-party libraries alone.

2.1 `std::math`

Specifically, this proposal introduces a new namespace `std::math` to clearly separate the new customizable contract from the current functions inherited from C.

- Math functions in this namespace allow customization through a member function or ADL.
- Their naming strives to be consistent with modern standard library convention.
- They leverage overloading and generic programming instead of prefix and suffix.

In addition, tools offering autocompletion would also benefit from a focused, dedicated namespace narrowing the relevant options early by category rather than name only.

The introduction of a new namespace and a new contract also provides an opportunity to resolve longstanding ambiguities such as the semantics of `abs` and `norm`. The specific meaning of math functions becomes especially consequential once user-defined customization is possible.

Finally, the combination of a reduced number of function names, more predictable naming patterns, better autocompletion, and clearer semantic of the operations should help improve the teachability issues mentioned previously.

2.1.1 Proposed naming

2.1.1.1 Expanded names

It seems that the modern Standard Library naming patterns prefer longer more-descriptive names. In that spirit, some operations have their name expanded, especially when their current denomination in `<cmath>` is not a universally accepted shorthand. For instance `std::cbrt` would be renamed as `std::math::cube_root`. Similarly for `sqrt`, `pow`, or `hypot`.

2.1.1.2 Kept names

Where the existing name is considered common-enough and unambiguous, it is kept. This is for instance the case for `sin`, `cos`, `abs`, `exp`, `log`.

2.1.1.3 Remapping

With `std::complex`, `std::abs` has been used to return the Euclidian norm, and `std::norm` to return the “field norm”. Considering the “field norm” is not a norm in the strict analytic or geometric sense, it is considered ambiguous.

To desambiguate this situation, these functions change their contract in `std::math`:

- `std::math::abs` returns the component-wise absolute value.
- `std::math::norm` is templated to the specificity type of norm it implements (L_1, L_2, L_∞), with a default to L_2 norm.
- `std::math::norm_squared` is defined and returns the square of the L_2 norm.

2.2 Bridging through opt-in

In addition to the new namespace, this proposal also introduces a mechanism in the `std` namespace to enable custom implementation for types explicitly opted-in, with math functions. This is a compatibility layer with a strict contract to maintain the existing code behaviour: custom implementations are only called for explicitly opted-in types and when the alternative resolution is failing compilation (ill-formed).

We acknowledge the limit that existing behaviours relying on SFINAE based on mathematical functions not being resolved and applied on explicitly opted-in types might break. Opt-in should thus be exercised with caution, especially in contexts where SFINAE is conditioned on mathematical function availability.

The behaviour of this compatibility layer is deliberately different from the behaviour in `std::math`: the former prioritizes backward compatibility with explicit opt-in and user-specified implementations as fallbacks; the latter prioritizes user-defined customization, without opt-in, with `std` variant as the fall-back.

The compatibility layer contract is deliberately narrow, opening a path where failure was the alternative. This is to prevent edge-cases where the opt-in could result in behaviour change at a distance. It leaves open the possibility of widening this contract in the future based on real usage data.

2.3 Scope

Although the introduction of a new namespace provides an opportunity to fix some contract ambiguity, the scope of this proposal is limited to enabling the extension of existing mathematical operations to user-defined types.

This involves using a different name for some operations in the new namespace (e.g.: to match the naming patterns of the standard library), or in some cases to define distinct names for an existing operation (e.g.: `std::norm` split in the new namespace into `std::math::norm<>` and `std::math::norm_squared`).

However, this proposal does **not** attempt to introduce other new features such as separate new mechanisms for selecting local rounding behaviours or performance/precision trade-offs. It limits itself with not precluding such additions to the language.

This proposal is intended to introduce a minimal extensible library for common math functions. It is composed of three main parts:

- the API of the new library,
- the compatibility layer in `std`,
- a minimal set of implementations for the built-in types.

This minimality is especially motivated by the precautionary principle and the concern for standard library implementers bandwidth.

2.3.1 Future extensions

Although it would be possible to use the dispatch mechanism to define specialised functions that fallback on default implementation if a specialised implementation is not found for the given type (e.g. a fast square root function dispatching to dedicated fast algorithm if it is found for the type of the argument, or defaulting to `std::math::square_root` otherwise), this is not leveraged as part of the current proposal. Users are free to make use of this possibility, but the provided implementations don't.

It does not provide algorithms or other mathematical functions leveraging other (potentially user-defined) mathematical operations it defines in their implementation. For instance, it would be possible to define the power operation in terms of exponentiation and logarithm by default. This possibility is left open for user-defined implementation but not provided as part of this proposal.

2.3.1.1 Rounding modes, precision/performance control

As stated before, adding such additional features to the language is out of the scope of this proposal. This said, if such additions were made simultaneously or in a future version of the language, the CPOs of `std::math` namespace could adapt to support it.

While it is difficult to know the form these new features would take, we can consider four options they could be implemented as:

- An additional template parameter.
- Square brackets.
- An additional parameter.
- New prefixed or suffixed functions.

In all of these cases, the CPOs in the `std::math` namespace could follow the same pattern in the same way. In the case of additional prefixed or suffixed functions, it could be decided whether the rationale for this choice in the namespace `std` holds for `std::math` and make a decision on this basis (considering if an additional template parameter would be preferred to maintain the design goal of minimizing the number of variants for the same operation).

2.4 Summary

This proposal doesn't require any new language feature. It would introduce a new `std::math` namespace and a compatibility layer inside of the `std` namespace, each of these two additions introducing approximately one new declaration per mathematical function, comparable in surface area to `<cmath>`.

3 Design Principles

3.1 Two spaces, different behaviours

The design space for this proposal is effectively split in two, between the forward-looking `std::math` layer, and the compatibility layer in `std`.

A solution in `std` only would have to compromise between extensibility and backward compatibility. Breaking the established contract there would be dangerous, and an opt-in-only solution leads to no good outcome: if adoption is low, generic code cannot rely on it for extensibility; if most types opt in, the mechanism ceases to be optional in any meaningful sense and only gets in the way.

Conversely, a new namespace alone would not help with code that calls `std` mathematical functions directly.

By combining both, the new code can benefit from the new approach in the new namespace with a new contract, while the compatibility layer in `std` remains available for targeted, case-by-case opt-ins.

The new `std::math` namespace splits from the mathematical functions inherited from C. It follows standard library naming patterns and leverages overloading and templates. It is intended as a native C++ library.

Compromising between C legacy and modern C++ patterns would hinder both objectives. The clear separation allows both sides to remain, each in its defined space.

3.2 New contract

The math functions in the `std::math` namespace have a new contract, independent from the `std` namespace contract.

User-customization implies a looser contract for these functions when they are dispatched to a user-defined implementation, where few aspects of such a contract can be enforced. Similarly to overloaded operators or other CPOs, the contract is owned by the implementation the call gets dispatched to.

The implementations for native types can define their own contracts, potentially forwarding the same contract as their `std` equivalent, but no such guarantee can be provided for user-defined implementations.

Since the new namespace adopts the standard library naming patterns, departing from the `<cmath>` names, several operations will have a different name between the two namespaces. In such cases, they will read as distinct functions, and the users will expect a different contract.

More importantly, with user-customization enabled, the function contract needs to be as clear as possible to avoid confusion. To that end, function names are chosen to reflect their purpose as precisely as possible.

Such an example of ambiguity arose with the `abs` and `norm` functions for the `std::complex` type, where `std::abs` was defined to be the L2 norm, and `std::norm` to be the square of the L2 norm (the squared modulus, which is not a mathematical norm). This would make it unclear how the `abs` and `norm` functions should be customized for user-defined multi-dimensional types.

3.3 CPO-like customization

Rather than a plain function template, `std::math` operations are proposed as CPO-like constructs. This design, established by the Ranges library (`std::ranges::begin`, `std::ranges::swap` etc.), offers these two advantages:

- The customization point cannot be found by ADL on user types, since it is an object rather than a function.
- The `operator()` can be constrained, making the customization point SFINAE-friendly and allowing it to be used correctly inside `requires` expressions and concepts.

3.3.1 Dispatch

The dispatch priority of these CPO-like constructs is explicitly:

1. To a member function (not subject to ADL).
2. To a free function found via ADL in the type's own namespace
3. To an equivalent `std` operation as the final fallback for all legacy types

The parameters are forwarded by forward reference.

User-defined implementations are not required to be `constexpr`, but the dispatch mechanism does not inhibit `constexpr` when it is supported.

3.3.2 Alternative to plain CPO

The term “CPO-like” is used instead of “CPO” directly to reflect the possibility of using upcoming improvement over the current CPO pattern. Specifically, if [P2806R3] (do expressions) and [P2826R2] (replacement functions) are accepted for C++29, the implementation of `std::math::square_root` reduces significantly.

The do expression makes the ADL idiom available in expression contexts, and replacement functions provide expression-equivalence guarantees. This would allow the entire CPO to be expressed as a single declaration, without explicit dispatch concepts or a poison pill. The direction proposed here would compose naturally with these future language improvements.

3.4 Single header

All functions are defined in a single header. The specific name of this header remains to be decided.

- `<math>` is easy to remember and consistent with other header names, but might conflict with `<math.h>`.
- `<cpmath>` or `<stdmath>` would avoid that confusion.

Specialisations for built-in types would also be defined in that same header. Specialisation for other types would be defined as part of the header associated with this type (e.g. `<complex>` for `std::complex`).

3.5 Canonical function names

Functions in the `std::math` namespace should define a single name per operation. Introducing name variants (prefix, suffix) should be avoided. Overloading and templating are preferred instead.

This helps keeping the surface area of the library as compact as possible and allows users to reach for the mathematical operation they need in a unified way.

The purpose of each mathematical operation should be made as clear as possible from its name and signature (as exercised by the native types) to make their user-customization as unambiguous as possible.

For the `abs` example, this would mean a single `abs` operation (no `labs`, `llabs`, `imaxabs`, `fabs`, `fabsf`, or `fabsl`) which is meant to return a value of the same type as its received parameter, where all negative components of the input value have had their sign switched (component-wise abs).

`norm` would be optionally templated on an enum defining which norm is requested (L1, L2, Infinity/Chebyshev), and defaulted to L2 when omitted. An additional `norm_squared` function could be considered, returning specifically the square of the L2 norm to offer a more performant alternative where the squared value is sufficient, without sacrificing the semantic of the operation.

See [Proposed naming](#).

3.6 Customisation-level specifications

The complexity, performance profile, rounding-mode and precision of an operation is typically defined at the customisation-level, allowing for types specialised for fast or precise implementations, or implementing a specific rounding mode.

For built-in types, the contract from the `std` version of the equivalent operation is usually maintained for backward-compatibility reasons.

4 General Specifications for built-in arithmetic types

Implementations for the fundamental arithmetic types (`char`, `short`, `int`, `long`, `long long`, `float`, `double`, `long double`, and `bool` where relevant) are provided where conservatively appropriate.

Specializations for `std::complex<T>` are provided where appropriate, with semantics defined independently from their `std` counterparts. All other standard library and user-defined types participate through the ADL dispatch mechanism.

These Standard Library implementations abide by the following rules.

4.1 Conservative addition

`std::math` functions are not specialized for every fundamental arithmetic type through this proposal. Specifically, when the contract of a specialization is unclear and its presence not indispensable, it is deferred to a later revision.

For instance, `std::math::square_root` return type when called on an integral type is unclear. It could be a `float`, a `double`, an `int`, or alternatively, this specialization of the operation could be templated to require an explicit return type as a template parameter. This proposal therefore chooses to leave `std::math::square_root` undefined for integral types, keeping this design space open for future improvements. This is both a precautionary principle and to avoid overwhelming implementers with less valuable work.

4.2 constexpr

All Standard-Library-defined specializations of functions in `std::math` are `constexpr`.

4.3 [[nodiscard]]

All functions in `std::math` with a return type are `[[nodiscard]]` unless specified otherwise.

4.4 Error Handling

Domain errors are considered contract violations and do not constitute runtime errors. As such they don't prevent the function from being `noexcept`.

`std::math` functions define their contract explicitly. They neither `throw` nor update `errno`.

Floating-point domain errors follow IEEE754 semantics (NaN/inf). Integral-type domain errors are Undefined Behaviour.

Standard Library specialisations for fundamental arithmetic types and `std::complex` are `noexcept`.

5 Technical Specification

5.1 Functions

5.1.1 Power Functions

5.1.1.1 `std::math::square_root`

Signature	std Equivalent
<code>R square_root(T x)</code>	<code>std::sqrt</code>

See implementation example.

5.1.1.1.1 Contract

Returns the square-root of the `val` parameter such that `square_root(x) * square_root(x) ≈ x`. The result's approximation can depend on the representable values allowed by the type, the rounding behaviour, and trade-offs between precision and performance.

The return type may differ from the input type ($\sqrt{m^2} \rightarrow m$).

5.1.1.1.2 Specializations

```
/*floating-point-type*/ square_root(/*floating-point-type*/)
```

Same contract as `std::sqrt(/*floating-point-type*/)`.

```
template<typename T>  
std::complex<T> square_root(const std::complex<T>&)
```

Same contract as `std::sqrt(const std::complex<T>&)`
(complex square root in the range of the right half-plane).

5.1.1.2 `std::math::cube_root`

Signature	std Equivalent
<code>R cube_root(T x)</code>	<code>std::cbrt</code>

5.1.1.2.1 Contract

Returns the cube-root of the `val` parameter such that `cube_root(x) * cube_root(x) * cube_root(x) ≈ x`. The result's approximation can depend on the representable values allowed by the type, the rounding behaviour, and trade-offs between precision and performance.

The return type may differ from the input type ($\sqrt[3]{m^3} \rightarrow m$).

5.1.1.2.2 Specializations

```
/*floating-point-type*/ cube_root(/*floating-point-type*/)
```

Same contract as `std::cbrt(/*floating-point-type*/)`

5.1.1.3 `std::math::power`

Signature	std Equivalent
R power(T base, U exp)	std::pow

5.1.1.3.1 Contract

Compute the value of `base` raised to the power `exp`.

5.1.1.3.2 Specializations

```
/* floating-point-type */
power ( /* floating-point-type */ base, /* floating-point-type */ exp )
```

Same contract as `std::pow(/*floating-point-type*/, /*floating-point-type*/)`.

```
template< class T1, class T2 >
std::complex<std::common_type_t<T1, T2>>
    power( const std::complex<T1>& x, const std::complex<T2>& y );

template< class T, class NonComplex >
std::complex<std::common_type_t<T, NonComplex>>
    power( const std::complex<T>& x, const NonComplex& y );

template< class T, class NonComplex >
std::complex<std::common_type_t<T, NonComplex>>
    power( const NonComplex& x, const std::complex<T>& y );
```

Same contracts as:

- `std::pow(const std::complex<T1>&, const std::complex<T2>&)`,
- `std::pow(const std::complex<T>&, const NonComplex&)`,
- `std::pow(const NonComplex&, const std::complex<T>&)`

respectively.

5.1.1.4 std::math::hypotenuuse

Signature	std Equivalent
T hypotenuuse(T x, T y, Ts...)	std::hypot ¹

(1) `std::hypot` is only equivalent for 2 to 3 arguments.

5.1.1.4.1 Contract

Returns $\sqrt{\sum x_i^2}$, computed so as to avoid intermediate overflow and underflow.

5.1.1.4.2 Specializations

```
template<class T, class... Ts>
requires (std::floating_point<T> && (std::same_as<T, Ts> && ...))
T hypotenuuse(T x, T y, Ts...);
```

Same contract as `std::hypot(/*floating-point-type*/, /*floating-point-type*/)` and `std::hypot(/*floating-p` with 2 and 3 arguments.

This same contract is generalised for more arguments.

5.1.2 Exponential

5.1.3 Trigonometric

5.1.4 Hyperbolic

5.1.5 Rounding

5.1.6 Sign

5.2 Proposed Implementation

This section presents some code illustrations implementing `std::math::square_root` as a representative example.

5.2.1 CPO

```
namespace std::math {
namespace __square_root {

    // Poison pill: prevents unqualified ADL calls from accidentally
    // finding std::math::square_root during concept checking
    void square_root(auto) = delete;

    template<typename T>
    concept has_member_square_root = requires(T&& x)
    {
        std::forward<T>(x).square_root();
    };

    template<typename T>
    concept has_adl_square_root = !has_member_square_root<T> && requires(T&& x)
    {
        square_root(std::forward<T>(x)); // ADL only; poison pill blocks std::math::square_root
    };

    template<typename T>
    concept has_std_sqrt = !has_member_square_root<T> && !has_adl_square_root<T> && requires(T&& x)
    {
        std::sqrt(std::forward<T>(x));
    };

    struct __fn
    {
        template<typename T>
        requires has_member_square_root<T>
            || has_adl_square_root<T>
            || has_std_sqrt<T>
        [[nodiscard]] constexpr auto operator()(T&& x) const
            noexcept(
                (has_member_square_root<T> && noexcept(std::forward<T>(x).square_root()))
                || (has_adl_square_root<T> && noexcept(square_root(std::forward<T>(x))))
                || (has_std_sqrt<T> && noexcept(std::sqrt(std::forward<T>(x))))
            )
        {
            if constexpr (has_member_square_root<T>)
                return std::forward<T>(x).square_root(); // member customization
            else if constexpr (has_adl_square_root<T>)
                return square_root(std::forward<T>(x)); // ADL
            else
                return std::sqrt(std::forward<T>(x)); // std fallback
        }
    };
}
```

```

    };
};

} // namespace __square_root

inline namespace __cpo {
    inline constexpr __square_root::__fn square_root{};
}
} // namespace std::math

```

5.2.2 Compatibility layer

```

namespace std {

// Opt-in trait for the compatibility forwarding layer.
// Users specialise this for their own types.
template<typename T>
inline constexpr bool is_math_extensible = false;

namespace __sqrt_compat {

    template<typename T>
    concept has_native_sqrt = requires(T&& x)
    {
        sqrt(std::forward<T>(x)); // resolves to the existing std::sqrt overload set,
                                // including built-in conversions
    };

} // namespace __sqrt_compat

// Forward to std::math for opted-in types, but only when no
// existing std::sqrt overload (including via conversion) applies.
template<typename T>
    requires is_math_extensible<std::remove_cvref_t<T>>
        && (!__sqrt_compat::has_native_sqrt<T>)
constexpr auto sqrt(T&& x)
    noexcept(noexcept(math::square_root(std::forward<T>(x))))
    -> decltype(math::square_root(std::forward<T>(x)))
{
    return math::square_root(std::forward<T>(x));
}

} // namespace std

```

5.2.3 Alternative implementation with [\[P2806R3\]](#) (do expressions) and [\[P2826R2\]](#) (replacement functions)

This example is for illustration purposes, and doesn't necessarily represent the proposal syntax.

```

namespace std::math
{
    static constexpr struct
    {
        using operator()(auto&& x) = (do { using std::sqrt; do_return sqrt(std::forward<decltype(x)>(x));
        } square_root;
    }
}

```


6 References

- [boost-units-sqrt] Boost Developers. Boost.Units: sqrt Implementation.
<https://github.com/boostorg/units/blob/develop/include/boost/units/cmath.hpp>
- [eigen-math] Eigen Developers. Eigen: MathFunctions.h Implementation.
<https://gitlab.com/libeigen/eigen/-/blob/master/Eigen/src/Core/MathFunctions.h>
- [mp-units] Mateusz Pusz. mp-units: Math Dispatch Implementation.
https://github.com/mpusz/mp-units/blob/b0e72810b983841b260d570b241c52586aa78999/src/core/include/mp-units/framework/representation_concepts.h#L227-L262
- [nholthaus-issue39] Nick Holthaus. nholthaus/units Issue #39: ADL and Math Functions.
<https://github.com/nholthaus/units/issues/39>
- [nholthaus-units] Nick Holthaus. nholthaus/units Library.
<https://github.com/nholthaus/units>
- [P2806R3] Barry Revzin, Bruno Cardoso Lopez, Zach Laine, Michael Park. 2025-01-12. do expressions.
<https://wg21.link/p2806r3>
- [P2826R2] Gašper Ažman. 2024-03-18. Replacement functions.
<https://wg21.link/p2826r2>
- [P3045R8] Mateusz Pusz, Dominik Berner, Johel Ernesto Guerrero Peña, Charles Hogg, Nicolas Holthaus, Roth Michaels, Vincent Reverdy. 2026-05-12. Quantities and units library.
<https://wg21.link/p3045r8>
- [P3935R1] Jan Schultke. 2026-05-11. Rebasing <cmath> on C23.
<https://wg21.link/p3935r1>
- [pusz-cpponseas2025] Mateusz Pusz. The 10 Essential Features for the Future of C++ Libraries. C++ on Sea 2025.
<https://www.youtube.com/watch?v=TJg37Sh9j78&t=3158s>